

Optimal Stochastic Trace Estimation

Christopher Musco

New York University, Tandon School of Engineering

COLLABORATORS



Raphael Meyer
(NYU)



Cameron Musco
(UMass. Amherst)



David Woodruff
(CMU)

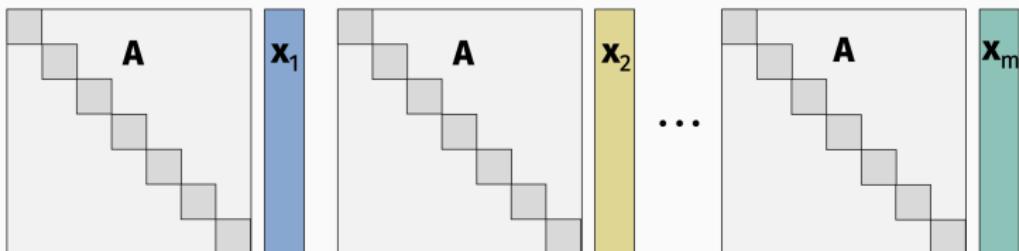
Paper available at: <https://arxiv.org/pdf/2010.09649.pdf>.

Recently accepted to the Symposium on Simplicity in
Algorithms (SOSA 2021).

IMPLICIT TRACE ESTIMATION

Basic problem in linear algebra:

- Given access to a $n \times n$ matrix A through a **matrix-vector multiplication oracle**.
- Goal is to (approximately) compute $\text{tr}(A) = \sum_{i=1}^n A_{ii}$.



Main question: How many matrix-vector multiplication “queries” Ax_1, \dots, Ax_m are required to compute $\text{tr}(A)$?¹

¹ x_i can be chosen adaptively, based on result of Ax_1, \dots, Ax_{i-1} .

IMPLICIT TRACE ESTIMATION

Algorithms in this model are called matrix-free,
or implicit matrix methods.

Typically useful when \mathbf{A} is not stored explicitly, but we have an efficient algorithm for multiplying \mathbf{A} by a vector.

Example: Hessian matrix-vector products.

Suppose we have some function $f(\mathbf{y})$ and we can efficiently compute gradients $\nabla f(\mathbf{y})$ for any \mathbf{y} . Let $\mathbf{A} = \nabla^2 f(\mathbf{y})$. Then:

$$\mathbf{Ax} \approx \frac{\nabla f(\mathbf{y} + \eta \mathbf{x}) - \nabla f(\mathbf{y})}{\eta} \quad \text{for sufficiently small } \eta.$$

IMPLICIT TRACE ESTIMATION

Also important when \mathbf{A} is a function of another matrix \mathbf{B} :

$$\mathbf{A} = f(\mathbf{B})$$

Common examples:

$$\mathbf{A} = \mathbf{B}^T \mathbf{B}$$

$$\mathbf{A} = \mathbf{B}^3$$

$$\mathbf{A} = 2\mathbf{B}^3 - 3\mathbf{B}^2 - \mathbf{I}$$

Cost to compute \mathbf{A} and $\text{tr}(\mathbf{A})$ explicitly:

$$O(n^3)$$

$$O(n^3)$$

$$O(n^3)$$

Cost to compute matrix-vector multiplication \mathbf{Ax} :

$$O(n^2)$$

$$O(n^2)$$

$$O(n^2)$$

All cheaper by a factor of $n!$ Even more savings if \mathbf{A} is sparse or structured.

IMPLICIT TRACE ESTIMATION

For more complex matrix functions, we can often compute $\mathbf{A}\mathbf{x} = f(\mathbf{B})\mathbf{x}$ efficiently using iterative methods:

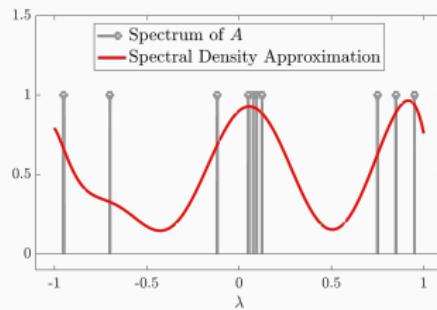
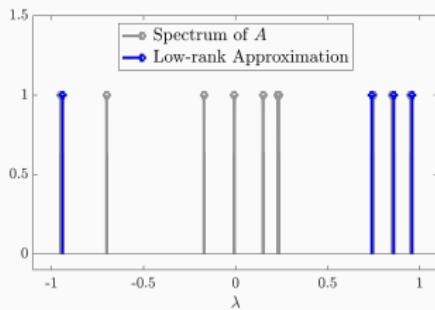
- Conjugate gradient, or any other system solver: $\mathbf{A} = \mathbf{B}^{-1}$.
- Lanczos method: $\mathbf{A} = \exp(\mathbf{B})$, $\mathbf{A} = \sqrt{\mathbf{B}}$, $\mathbf{A} = \log(\mathbf{B})$, etc.

All run in $O(n^2 \cdot C)$ time, where C depends on properties of \mathbf{B} .
For example, for $\mathbf{A} = \mathbf{B}^{-1}$, $C = \sqrt{\kappa} \cdot \log(1/\epsilon)$.

In practice, we typically have $O(n^2 \cdot C) \ll O(n^3)$.

EXAMPLE APPLICATIONS

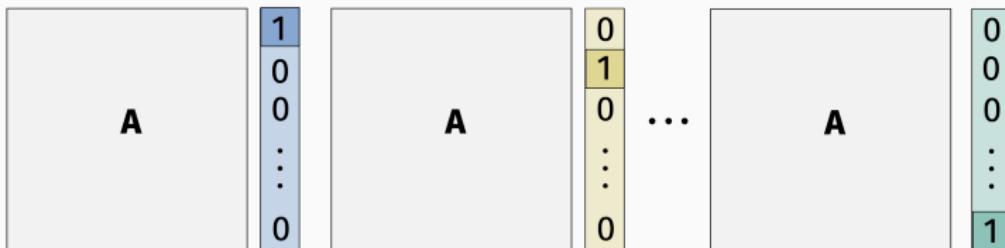
- Log-likelihood computation in Bayesian optimization, experimental design. $\text{tr}(\log(\mathbf{B})) = \text{logdet}(\mathbf{B})$.
- Estrada index, network connectivity. $\text{tr}(\exp(\mathbf{B}))$.
- Triangle counting in graphs. $\text{tr}(\exp(\mathbf{B}^3))$.
- Counting number of eigenvalues in an interval.
- Spectral density estimation.
- Matrix norms.



NAIVE EXACT ALGORITHM

Naive approach:

- Set $\mathbf{x}_i = \mathbf{e}_i$ for $i = 1, \dots, n$.
- Return $\text{tr}(\mathbf{A}) = \sum_{i=1}^n \mathbf{x}_i^T \mathbf{A} \mathbf{x}_i$



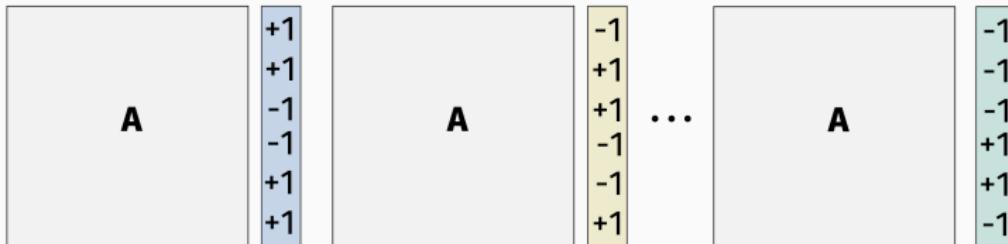
Returns exact solution, but requires n matrix-vector multiplies.
We want $\ll n$ multiplies, and will do so by allowing for
approximation.

HUTCHINSON'S STOCHASTIC TRACE ESTIMATOR

Simple, powerful, and widely used method for trace estimation.

Hutchinson 1991, Girard 1987:

- Draw $\mathbf{x}_1, \dots, \mathbf{x}_m \in \mathbb{R}^n$ i.i.d. with random $\{+1, -1\}$ entries.
- Return $\tilde{T} = \frac{1}{m} \sum_{i=1}^m \mathbf{x}_i^T \mathbf{A} \mathbf{x}_i$ as approximation to $\text{tr}(\mathbf{A})$.



HUTCHINSON'S STOCHASTIC TRACE ESTIMATOR

Let \tilde{T} be the trace estimate returned by Hutchinson's method.

Claim (Avron, Toledo 2011, Roosta, Ascher 2015)

If $m = O\left(\frac{\log(1/\delta)}{\epsilon^2}\right)$, then with probability $(1 - \delta)$,

$$\left| \tilde{T} - \text{tr}(A) \right| \leq \epsilon \|A\|_F.$$

If A is symmetric positive semidefinite (PSD) with eigenvalues $\lambda_1, \dots, \lambda_n$, then

$$\|A\|_F = \sqrt{\sum_{i=1}^n \lambda_i^2} \leq \sum_{i=1}^n \lambda_i = \text{tr}(A).$$

Corollary: For PSD A : $(1 - \epsilon) \text{tr}(A) \leq \tilde{T} \leq (1 + \epsilon) \text{tr}(A)$.

EXPECTED VALUE ANALYSIS

Hutchinson's Estimator:

- Draw $\mathbf{x}_1, \dots, \mathbf{x}_m \in \mathbb{R}^n$ i.i.d. with random $\{+1, -1\}$ entries.
 - Return $\tilde{T} = \frac{1}{m} \sum_{i=1}^m \mathbf{x}_i^T \mathbf{A} \mathbf{x}_i$ as approximation to $\text{tr}(\mathbf{A})$.
-

Expected value analysis:

For a single random ± 1 vector \mathbf{x} ,

$$\mathbb{E}[\tilde{T}] = \mathbb{E}[\mathbf{x}^T \mathbf{A} \mathbf{x}] = \mathbb{E} \sum_{i=1}^n \sum_{j=1}^n x_i x_j \mathbf{A}_{ij} = \sum_{i=1}^n \sum_{j=1}^n \mathbb{E}[x_i x_j \mathbf{A}_{ij}] = \sum_{i=1}^n \mathbf{A}_{ii}$$

So the estimator is correct in expectation:

$$\mathbb{E}[\tilde{T}] = \text{tr}(\mathbf{A}).$$

Hutchinson's Estimator:

- Draw $\mathbf{x}_1, \dots, \mathbf{x}_m \in \mathbb{R}^n$ i.i.d. with random $\{+1, -1\}$ entries.
 - Return $\tilde{T} = \frac{1}{m} \sum_{i=1}^m \mathbf{x}_i^T \mathbf{A} \mathbf{x}_i$ as approximation to $\text{tr}(\mathbf{A})$.
-

Variance analysis:

$$\begin{aligned}\text{Var}[\tilde{T}] &= \frac{1}{m} \text{Var}[\mathbf{x}^T \mathbf{A} \mathbf{x}] = \frac{1}{m} \text{Var} \left[\sum_{i=1}^n \sum_{j=1}^n x_i x_j A_{ij} \right] \\ &= \frac{1}{m} \sum_{i=1}^n \sum_{j=1}^n \text{Var}[x_i x_j A_{ij}] = \frac{1}{m} \sum_{i \neq j}^n A_{ij}^2 \leq \frac{1}{m} \|\mathbf{A}\|_F^2\end{aligned}$$

(We used that $x_i x_j$ and $x_j x_k$ are pairwise independent.)

FINAL ANALYSIS

Hutchinson's Estimator:

- Draw $\mathbf{x}_1, \dots, \mathbf{x}_m \in \mathbb{R}^n$ i.i.d. with random $\{+1, -1\}$ entries.
 - Return $\tilde{T} = \frac{1}{m} \sum_{i=1}^m \mathbf{x}_i^T \mathbf{A} \mathbf{x}_i$ as approximation to $\text{tr}(\mathbf{A})$.
-

Final analysis: Chebyshev's inequality implies that, with probability $9/10$,

$$|\tilde{T} - \text{tr}(\mathbf{A})| \leq \frac{1}{\sqrt{m/10}} \|\mathbf{A}\|_F.$$

Setting $m = O(1/\epsilon^2)$ gives $|\tilde{T} - \text{tr}(\mathbf{A})| \leq \epsilon \|\mathbf{A}\|_F$.

Getting correct $\log(1/\delta)$ dependence requires a bit more work
(Hanson-Wright inequality).

RESEARCH QUESTION

Result: $O(1/\epsilon^2)$ matrix-vector multiplies suffice to return, with prob. 9/10, a trace estimate for a PSD matrix with relative error:

$$(1 - \epsilon) \text{tr}(A) \leq \tilde{T} \leq (1 + \epsilon) \text{tr}(A).$$

Research Question: Is this tight?

RESEARCH QUESTION

Broader line of work: Tight upper bounds and lower bounds on complexity of basic linear algebra problems in “matrix-vector query” model.

- **Top eigenvector:** Simchowitz, Alaoui, Recht, 2018.
- **Least squares regression:** Braverman, Hazan, Simchowitz, Woodworth, 2020.
- **Rank, symmetry test, and more:** Sun, Woodruff, Yang, and Zhang, 2019.

MATRIX-VECTOR QUERY MODEL

The **matrix-vector query model** generalizes the most common models of computations in linear algebra.

Krylov subspace model:

- Compute $\mathbf{A}\mathbf{x}, \mathbf{A}^2\mathbf{x}, \dots, \mathbf{A}^m\mathbf{x}$ for chosen vector \mathbf{x} .
- Lower bounds typically via approximation theoretic arguments (understanding the limits of polynomials).

Matrix sketching model:

- Compute $\mathbf{A}\mathbf{x}_1, \dots, \mathbf{A}\mathbf{x}_m$ where $\mathbf{x}_1, \dots, \mathbf{x}_m$ are chosen non-adaptively (usually chosen to be random vectors).
- Lower bounds typically via one-round communication complexity.

Merits of this model:

- Captures most algorithms that are used in practice, where matrix-vector multiplies often dominate computation cost.
- Allowing arbitrary adaptivity makes the model quite a bit richer. Proving lower bounds seems harder but doable.
- Appears to be a “sweet spot” for understanding problem complexity in linear algebra.

Limitation:

- Does not capture methods like stochastic gradient or coordinate descent.

OUR RESULTS

Upper bound: $O(1/\epsilon)$ matrix-vector multiplies suffice to return, with prob. $9/10$, a trace estimate for a PSD matrix with relative error:

$$(1 - \epsilon) \text{tr}(\mathbf{A}) \leq \tilde{T} \leq (1 + \epsilon) \text{tr}(\mathbf{A}).$$

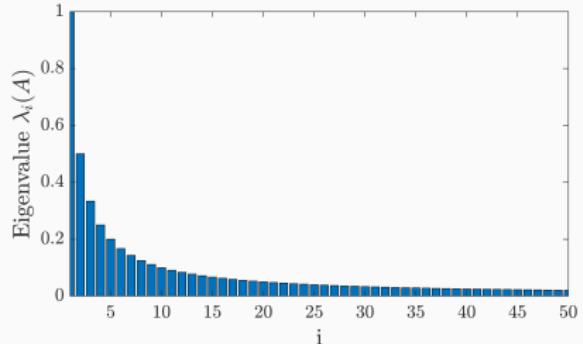
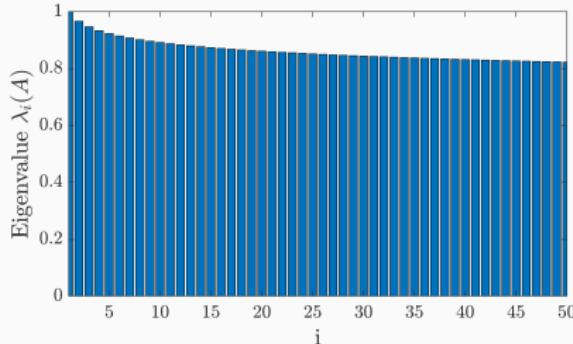
- Quadratic improvement over Hutchinson's $O(1/\epsilon^2)$.
- Algorithm achieving bound is nearly as simple.
- Performs much better experimentally.

Lower bound: $\Omega(1/\epsilon)$ matrix-vector multiplies are necessary to obtain a relative error approximation with probability $> 2/3$.

- Two different approaches: reduction from multi-round communication complexity, and from hypothesis testing for negatively spiked covariance matrices.

SPECTRUM DEPENDENT BOUND

Observation: Hutchinson's method performs much better when \mathbf{A} has a “flatter” spectrum.



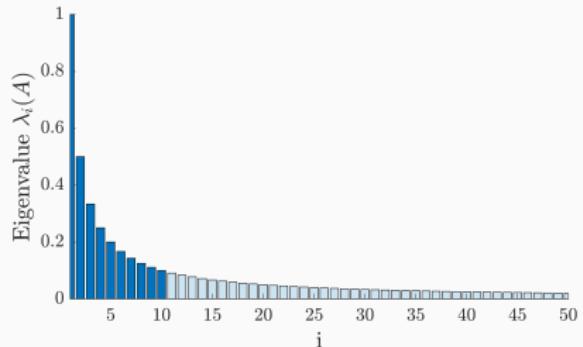
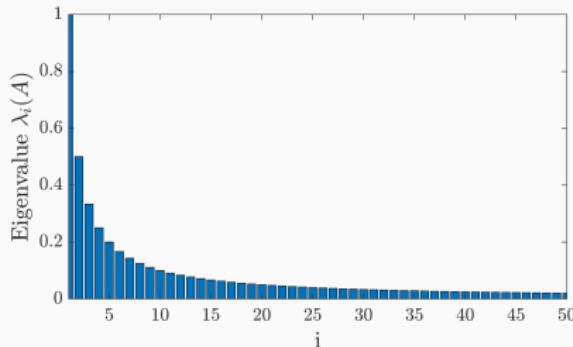
We proved that: $|\tilde{T} - \text{tr}(\mathbf{A})| \leq \epsilon \|\mathbf{A}\|_F \leq \epsilon \text{tr}(\mathbf{A})$, but when the spectrum is decaying $\|\mathbf{A}\|_F \ll \text{tr}(\mathbf{A})$.

In the extreme case when $\lambda_1 = \lambda_2 = \dots = \lambda_n$, we have:

$$\|\mathbf{A}\|_F = \sqrt{\sum_{i=1}^n \lambda_i^2} = \frac{1}{\sqrt{n}} \sum_{i=1}^n \lambda_i = \frac{1}{\sqrt{n}} \text{tr}(\mathbf{A}).$$

STEEP SPECTRUM

On the other hand, when \mathbf{A} 's spectrum is decaying, we get a good approximation by simply computing its top eigenvectors.



$$\text{tr}(\mathbf{A}) = \sum_{i=1}^n \lambda_i \approx \sum_{i=1}^k \lambda_i = \text{tr}(\mathbf{AQQ}^T)$$

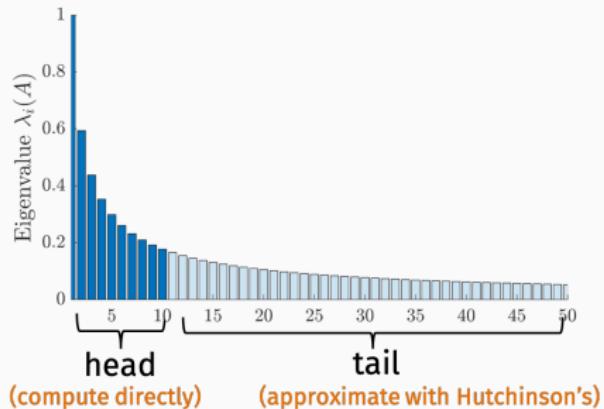
where $\mathbf{Q} \in \mathbb{R}^{n \times m}$ is an orthonormal span \mathbf{A} 's top k eigenvalues.

- \mathbf{Q} itself can be computed with $\sim O(k)$ matrix-vector multiplication queries using block power method or a Krlyov method (Saibaba, Alexanderian, Ipsen, 2018).
- Then $\text{tr}(\mathbf{A}\mathbf{Q}\mathbf{Q}^T) = \text{tr}(\mathbf{Q}^T(\mathbf{A}\mathbf{Q}))$ can be computed with k additional matrix-vector multiplies.

Main observation: Every spectrum is either “flat enough” or “decaying enough” to prove a better bound than $O(1/\epsilon^2)$.

OUR METHOD: HUTCH++

1. Find approximate span for top k eigenvectors \mathbf{Q} .
2. Observe that $\text{tr}(\mathbf{A}) = \text{tr}(\mathbf{AQ}\mathbf{Q}^T) + \text{tr}(\mathbf{A}(\mathbf{I} - \mathbf{QQ}^T))$
3. Approximate $\tilde{P} = \text{tr}(\mathbf{A}(\mathbf{I} - \mathbf{QQ}^T))$ using Hutchinson's with ℓ vectors.
4. Return $\tilde{T} = \text{tr}(\mathbf{AQ}\mathbf{Q}^T) + \tilde{P}$.



The only error is from the estimator for $\text{tr}(\mathbf{A}(\mathbf{I} - \mathbf{QQ}^T))$, which will have much lower variance if $\|\mathbf{A}(\mathbf{I} - \mathbf{QQ}^T)\|_F \ll \|\mathbf{A}\|_F$.

Standard result in Randomized Numerical Linear Algebra:

Lemma (Sarlos 2006, Woodruff 2014)

If $\mathbf{S} \in \mathbb{R}^{n \times m}$ is chosen with i.i.d. ± 1 entries, then $\mathbf{Q} = \text{orth}(\mathbf{AS})$ satisfies with probability $(1 - \delta)$,

$$\|\mathbf{A} - \mathbf{AQ}\mathbf{Q}^T\|_F \leq 2\|\mathbf{A} - \mathbf{A}_k\|_F,$$

as long as \mathbf{S} has $m = O(k + \log(1/\delta))$ columns.

Here \mathbf{A}_k is the best k -rank approximation to \mathbf{A} , obtained by projecting onto \mathbf{A} 's top k eigenvectors.

Note that \mathbf{Q} can be view as the result of running a single step of power method on \mathbf{A} .

FINAL BOUND

For any PSD matrix \mathbf{A} :

$$\|\mathbf{A} - \mathbf{A}_k\|_F^2 = \sum_{i=k+1}^n \lambda_i^2 \leq \lambda_{k+1} \sum_{i=k+1}^n \lambda_i \leq \frac{1}{k} \text{tr}(\mathbf{A}) \cdot \text{tr}(\mathbf{A}).$$

So if $\|\mathbf{A}(\mathbf{I} - \mathbf{Q}\mathbf{Q}^T)\|_F \leq 2 \|\mathbf{A} - \mathbf{A}_k\|_F$, then with high probability,

$$|\text{tr}(\mathbf{A}(\mathbf{I} - \mathbf{Q}\mathbf{Q}^T)) - \tilde{P}| \leq \frac{1}{\sqrt{\ell}} \|\mathbf{A}(\mathbf{I} - \mathbf{Q}\mathbf{Q}^T)\|_F \leq \frac{1}{\sqrt{\ell}} \cdot \frac{2}{\sqrt{k}} \text{tr}(\mathbf{A}).$$

Setting $\ell = k = O(1/\epsilon)$ gives error $\epsilon \text{tr}(\mathbf{A})$ and thus:

$$|\text{tr}(\mathbf{A}) - \tilde{T}| = |\text{tr}(\mathbf{A}(\mathbf{I} - \mathbf{Q}\mathbf{Q}^T)) - \tilde{P}| \leq \epsilon \text{tr}(\mathbf{A}).$$

FINAL ALGORITHM

Theorem (Final Result)

If $m = O\left(\frac{\log(1/\delta)}{\epsilon}\right)$ and A is PSD then with probability $(1 - \delta)$, Hutch++ returns \tilde{T} satisfying:

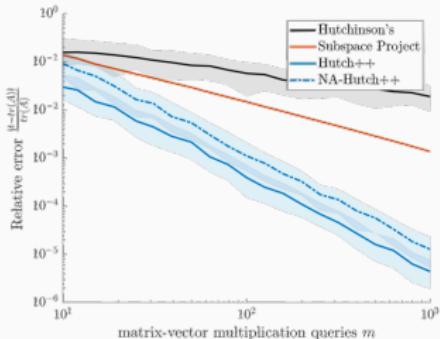
$$(1 - \epsilon) \operatorname{tr}(A) \leq \tilde{T} \leq (1 + \epsilon) \operatorname{tr}(A)$$

```
1  function T = hutchplusplus(A, m)
2 -
3 -     S = 2*randi(2,size(A,1),m/3);
4 -     G = 2*randi(2,size(A,1),m/3);
5 -     [Q,~] = qr(A*S,0);
6 -     G = G - Q*(Q'*G);
7 -     T = trace(Q'*A*Q) + 1/size(G,2)*trace(G'*A*G);
end
```

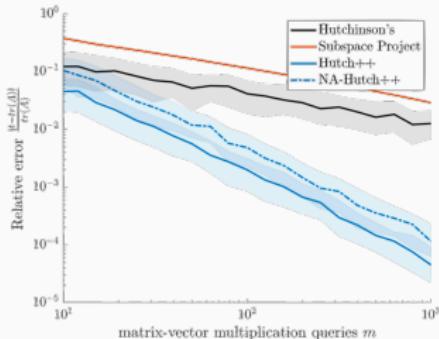
This algorithm is adaptive, meaning that the choice of x_i depends on Ax_1, \dots, Ax_{i-1} . We also have a non-adaptive method, NA-Hutch++ that achieves the same bound.

EXPERIMENTAL RESULTS

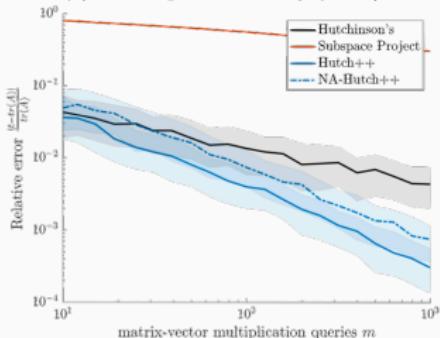
Results on synthetic matrix A with spectrum $\lambda_i = i^{-c}$ for different values of c .



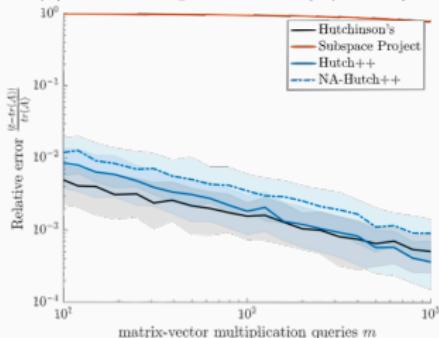
(a) Fast Eigenvalue Decay ($c = 2$)



(b) Medium Eigenvalue Decay ($c = 1.5$)



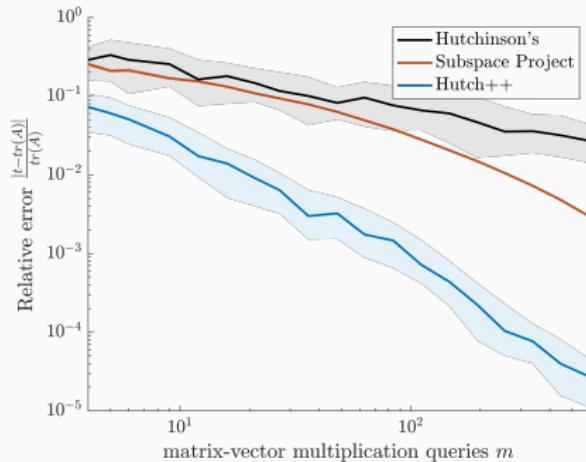
(c) Slow Eigenvalue Decay ($c = 1$)



(d) Very Slow Eigenvalue Decay ($c = .5$)

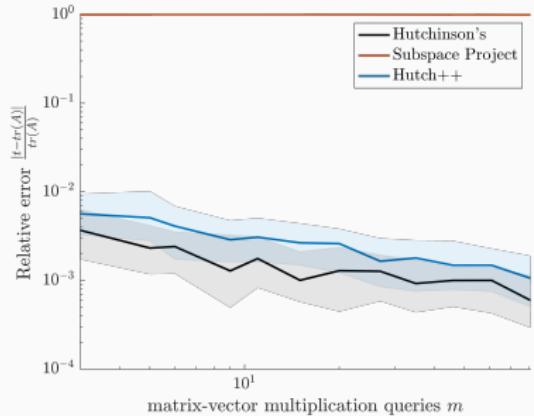
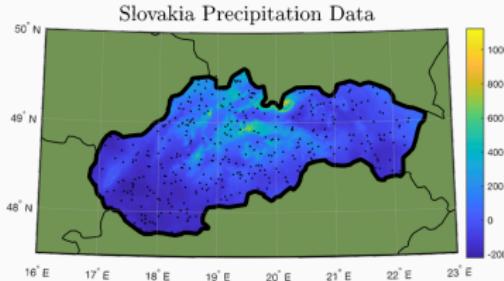
APPLICATIONS

If \mathbf{B} is symmetric with eigendeposition $\mathbf{V}\Lambda\mathbf{V}^T$, we let $f(\mathbf{B})$ denote $\mathbf{V}f(\Lambda)\mathbf{V}^T$, which means that f is applied entrywise to the diagonal matrix of eigenvalues, Λ . Note that $\text{tr}(\mathbf{B}) = \sum_{i=1}^n f(\lambda_i)$.



$\mathbf{A} = \exp(\mathbf{B})$ for graph adjacency matrix \mathbf{B} from linguistics application.
 $\text{tr}(\mathbf{A})$ is the well known Estrada Index or “natural connectivity”.

APPLICATIONS



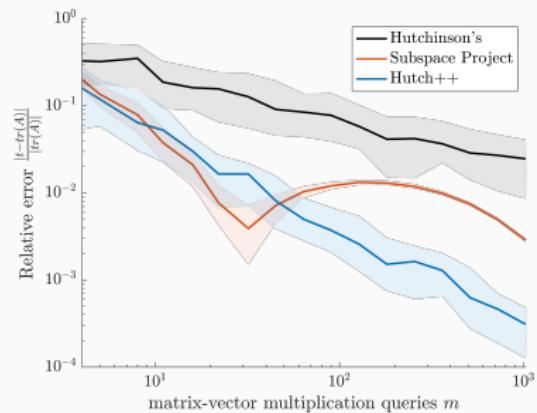
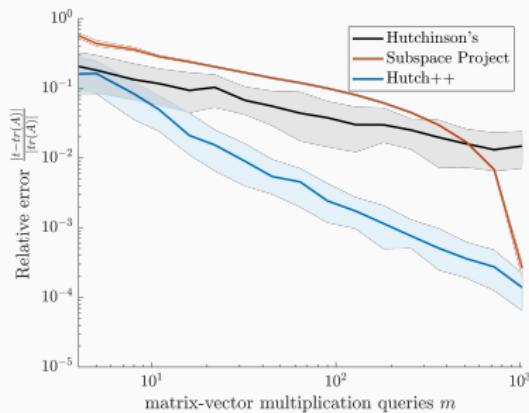
$A = \log(B + \lambda I)$ for kernel matrix B from Gaussian process regression.
 $\text{tr}(A) = \log \det(B)$, which is used in loglikelihood calculations.

Takeaway: For matrix functions that flatten B 's spectrum, Hutchinson's estimator performs far better than the $O(1/\epsilon^2)$ bound predicts. Hutch++ will never perform much worse.

APPLICATIONS

Hutch++ works well empirically for many non-PSD matrices.

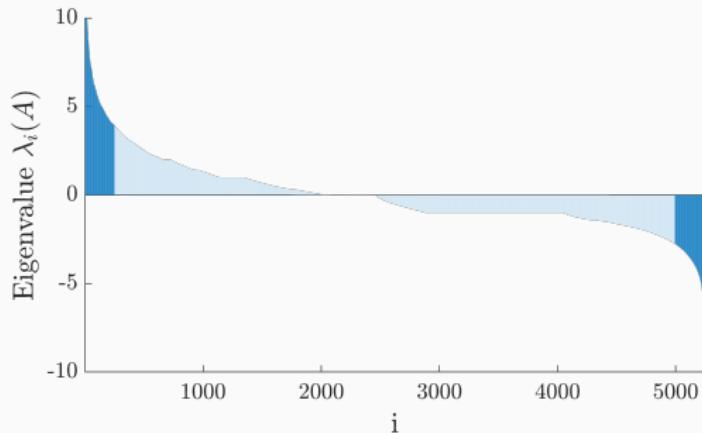
Let \mathbf{B} is the (indefinite) adjacency matrix of an undirected graph G , $\text{tr}(\mathbf{B}^3)$ is exactly equal to the number of triangles in G .



$A = B^3$ for arXiv.org citation network and Wikipedia voting network.

REAL APPLICATIONS

For non-PSD \mathbf{A} , the projection step, $\mathbf{A}(\mathbf{I} - \mathbf{Q}\mathbf{Q}^T)$ approximately removes \mathbf{A} 's largest magnitude eigenvalues, which can still reduce variance substantially.



Spectrum of $\mathbf{A} = \mathbf{B}^3$ for arXiv.org citation network.

Theorem

Any algorithm that accesses a PSD matrix \mathbf{A} via matrix-vector multiplication queries $\mathbf{Ax}_1, \dots, \mathbf{Ax}_m$, where $\mathbf{x}_1, \dots, \mathbf{x}_m$ are possibly adaptively chosen vectors with integer entries in $\{-2^b, \dots, 2^b\}$, needs

$$m = \Omega\left(\frac{1}{\epsilon \cdot [b + \log(1/\epsilon)]}\right) \text{ queries}$$

to approximate $\text{tr}(\mathbf{A})$ to multiplicative error $(1 \pm \epsilon)$.

Reduction to 2-party multi-round communication problem.
“Hard” input distribution will involve \mathbf{A} with integer entries, which is why we need the bit complexity bound b .

GAP HAMMING PROBLEM

Problem (Gap Hamming)

Let Alice and Bob be communicating parties who hold vectors $s, t \in \{-1, 1\}^n$, respectively. Must decide with few bits of communication if:

$$\langle s, t \rangle \geq \sqrt{n}$$

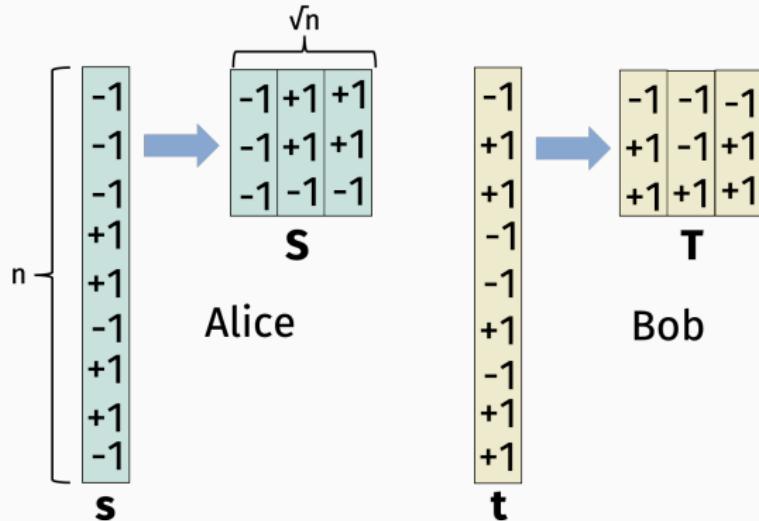
or

$$\langle s, t \rangle \leq -\sqrt{n}$$

Theorem (Chakrabarti, Regev 2012)

The randomized communication complexity for solving Problem 1 with probability $\geq 2/3$ is $\Omega(n)$ bits.

REDUCTION TO TRACE ESTIMATION



Let $Z = S + T$ and $A = Z^T Z$.

$$\text{tr}(A) = \|Z\|_F^2 = \|s + t\|_2^2 = 2n - 2\langle s, t \rangle.$$

So if Alice and Bob can estimate $\text{tr}(A)$ up to error $(1 \pm 1/\sqrt{n})$, then they will solve the Gap Hamming problem.

REDUCTION TO TRACE ESTIMATION

Claim: Alice and Bob can simulate any m query algorithm for estimating the trace of $\mathbf{A} = (\mathbf{S} + \mathbf{T})^T(\mathbf{S} + \mathbf{T})$ with $O(m\sqrt{n}(\log n + b))$ bits of communication.

- Alice decides on \mathbf{x}_1 , sends to Bob with $\sqrt{n} \cdot \log(2^b)$ bits.
- Bob computes $\mathbf{T}\mathbf{x}_1$, sends to Alice with $\sqrt{n} \cdot \log(\sqrt{n}2^b)$ bits.
- Alice computes $(\mathbf{S} + \mathbf{T})\mathbf{x}_1$.
- Repeat to multiply $(\mathbf{S} + \mathbf{T})\mathbf{x}_1$ by $(\mathbf{S} + \mathbf{T})^T$
- Alice decides on \mathbf{x}_2 , process repeats m times.

So, by $\Omega(n)$ lower bound for Gap Hamming, we can't have m less than $\frac{\sqrt{n}}{\log n + b}$. Setting $\epsilon = 1/\sqrt{n}$ gives the result.

OPEN QUESTIONS

- In progress: Lower bounds for e.g. $\text{tr}(A^3)$, $\text{tr}(\exp(A))$, $\text{tr}(A^{-1})$.
- What about (coarse) approximate matrix vector multiplications? We have some upcoming work on this related to spectral density estimation problems, but there's a lot to think about.
- Relates to model where we sample rows or columns of A (and implement things like SGD/SCD).
- Can we get conditional lower bounds for simple problems like triangle counting in a completely general computational model?

THANKS! QUESTIONS?